


Bidirectional Encoder (e.g., XLM-R)

Masked token prediction



Causal Decoder (e.g., Llama 3.1)

Next-token prediction

